# facebook

#### facebook

# Global Transaction ID at Facebook

Santosh Praneeth Banda, Evan Elias, and Yoshinori Matsunobu

# Agenda

- 1 Introduction to GTID
- 2 Benefits and use-cases
- 3 Failover, semi-sync, and MHA
- 4 FB MySQL Server improvements
- 5 Automation and monitoring
- 6 Deployment strategy

# Introduction to GTID

### **GTID Concepts**

- Traditional MySQL replication uses relative coordinates
  - Replication position is entirely relative to the current master
- With GTID a unique identifier is assigned to every transaction
  - Each instance tracks which GTIDs they have executed
  - Relative coordinates are also still available
- Each transaction with a given GTID will only be executed once per server

# MySQL 5.6 Implementation

- Each master has a UUID
- Each GTID looks like source\_uuid:transaction\_id
- transaction\_id is monotonically increasing per source
- Say server A has UUID 3E11FA47-71CA-11E1-9E33-C80AA9429562, then GTID 3E11FA47-71CA-11E1-9E33-C80AA9429562: 23 indicates the 23rd write transaction to originate here
- Transaction counter ordering is per source, not global!
  - If we then promote server B, its transaction IDs begin at 1
  - If we promote back to A, its transaction IDs resume at 24
  - Binlog is the source of truth for replaying transactions in order

### MySQL 5.6 GTID sets

- Servers track which range of GTIDs they've executed
  - gtid\_executed global variable
  - Also available in SHOW MASTER STATUS and SHOW SLAVE STATUS
  - Example: 3E11FA47-71CA-11E1-9E33-C80AA9429562:1-5
- · Server also keeps track of set of GTIDs in binlogs that have been purged

# MySQL 5.6 auto-positioning

- CHANGE MASTER TO ... MASTER\_AUTO\_POSITION = 1
- Omit MASTER\_LOG\_FILE and MASTER\_LOG\_POS
- Scans master's binlogs to find transactions that have not yet been executed on this replica
- If the replica is missing transactions that have already been purged from master's binlogs, I/O thread will error

# Other implementations

- Google's MySQL 5.0 GTID patch, by Justin Tolmer
  - Supported auto-positioning and crash-safe replication
  - Tracks latest GTID instead of using GTID sets
  - Didn't support master-master topologies
- MariaDB 10.x GTID support
  - Tracks latest GTID instead of using GTID sets
  - Uses a notion of "domain ID" to support master-master topologies
  - Permits online rollout
  - Commands and syntax all differ from Oracle's implementation

# Benefits and Use-cases

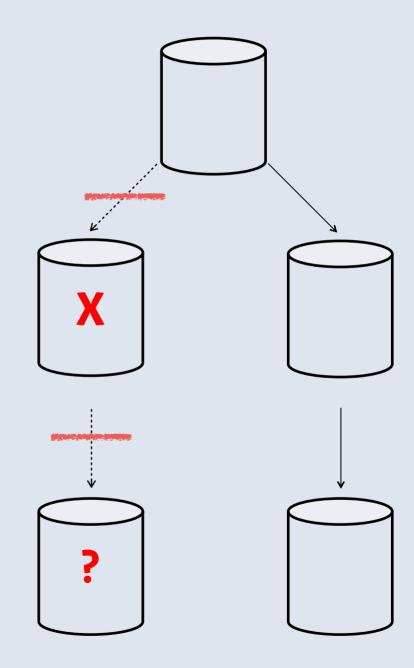
# Major areas of benefit

- Failover
  - Simpler, faster, less error-prone
  - Trivial to get slaves in sync after master failure
  - Server will safely ignore transactions it has already executed
- Backups
  - Cornerstone for point-in-time recovery from a single binlog stream
  - No need for duplicate streams per replica
- Hierarchical replication
  - Slaves-of-slaves are now much easier to manage

# Hierarchical replication chains

 Previously, if middle tier instance fails, no easy way to repoint its bottom-level replicas to another master

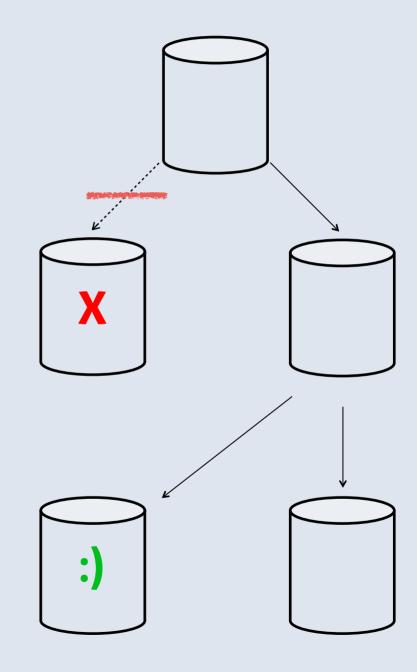
 Replication coordinates were always relative to the slave's immediate master!



# Hierarchical replication chains

 GTID auto-positioning fixes this.
 Trivial to repoint any member of the replica set to any other.

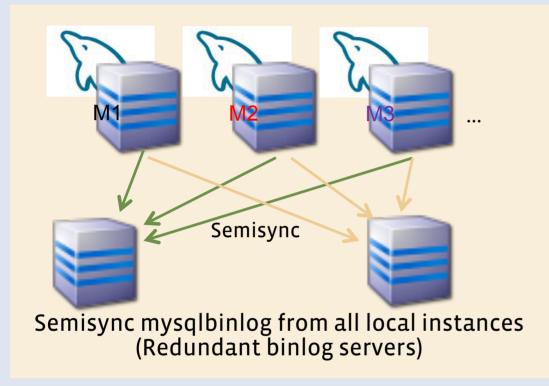
 Hierarchical replication not currently used at Facebook, but perhaps we may find a use-case in the future



# Failover, semi-sync, and MHA

# GTID and Loss-Less Semisync mysqlbinlog

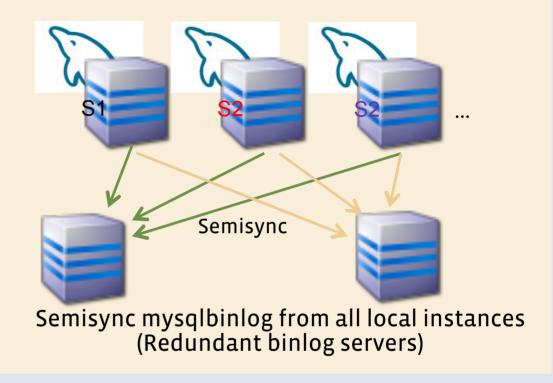
DC1 (primary master region)

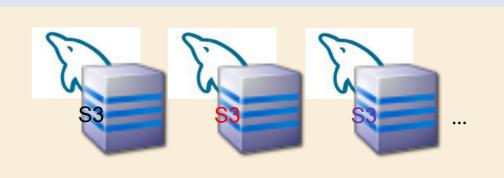


Async Repl

- Our mysqlbinlog speaks semisync replication protocol
- No dedicated local semisync slave needed
- Master failover is a bit tricky because semisync mysqlbinlog isn't a real mysqld

DC2 (secondary master region)

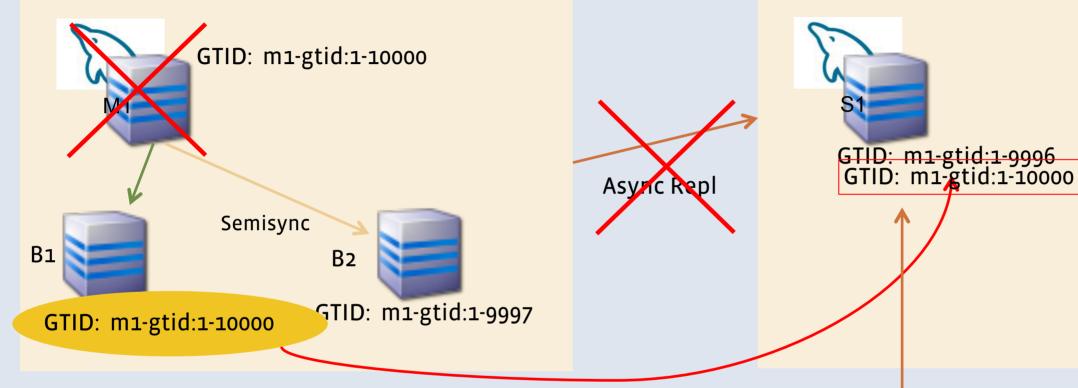




# Master failover with GTID + mysqlbinlog

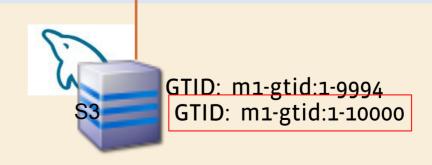
DC1 (primary master region)

DC2 (secondary master region)



#### Master Failover Steps:

- 1. Identify slave or binlog server with the latest GTID (B1)
- 2. If binlog server has the latest event, apply the diff GTID to a new master
- 3. Other slaves reconnect to the new master



# Extending MHA for GTID and binlog server

- MHA is a fast master failover tool, working with official MySQL 5.0~5.6, even without GTID
- We use semisync binlog servers, which don't accept MySQL commands
- Official tool mysqlfailover covers master failover with GTID, but it doesn't cover binlog servers
- We extended MHA to support both GTID and binlog servers
- https://code.google.com/p/mysql-master-ha/

### MHA and GTID failover

- MHA (Version 0.56) automatically detects whether to do GTID based failover
- GTID-based failover process:
  - Checking if binlog servers are available or not
  - Checking the latest GTID position
  - Recovering based on GTID (not using relay logs)

# Crash-safe master via loss-less semisync

#### What is crash-safe master?

- When master is down and promoting a slave, after the crashed master's recovery, we want to add the crashed master as a new slave without rebuilding the whole crashed master instance
- Recovery (recovering from OS reboot etc) shouldn't take days. Applying a few hours of binlogs is much faster than rebuilding entire instance



After crashed master's recovery, continue replication by CHANGE MASTER TO MASTER\_HOST='S1', MASTER\_AUTO\_POSITION=1;

### Extensions needed for crash-safe master

- Using InnoDB only
- Writing GTID to InnoDB (logfile or table) at transaction commit
- GTID and InnoDB must be consistent. Writing GTID to InnoDB logfile/ table guarantees this
- Using loss-less semisync (backporting from 5.7 semisync)
- Crashed master's binlog position must not be ahead any other slave (binlog reader). Loss-less semisync guarantees this.

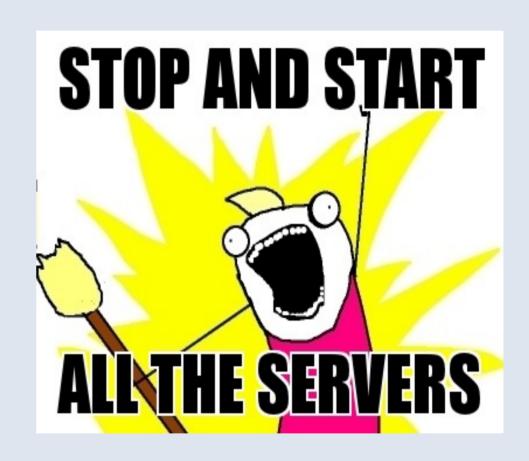
# FB MySQL Server improvements

# Facebook GTID improvements

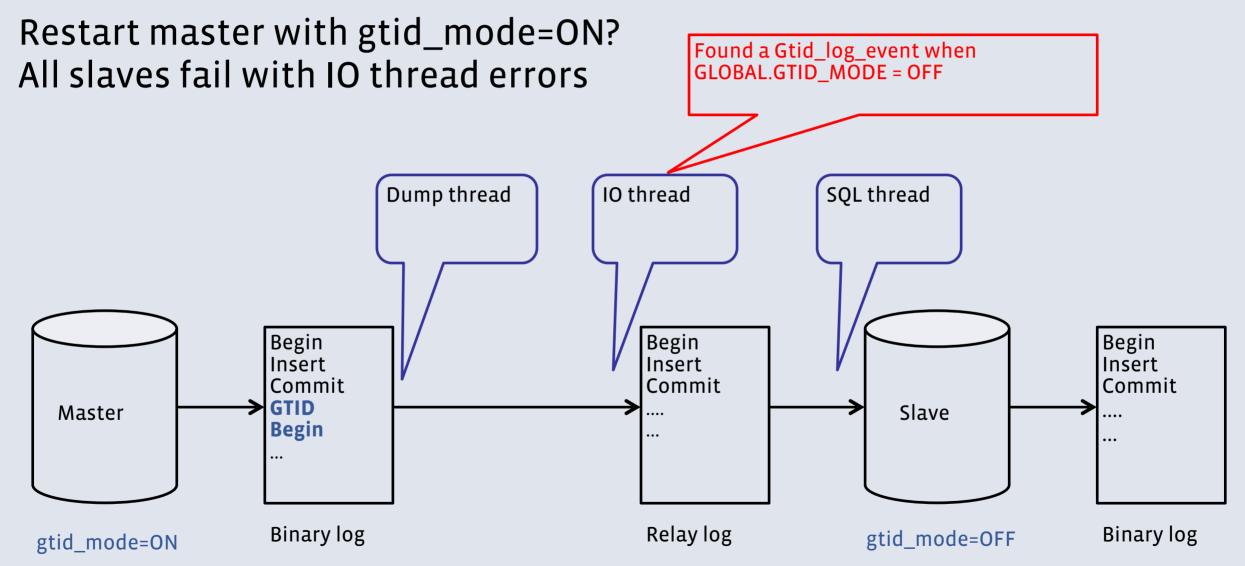
- Simplified deployment
- Restart and slave connect performance improvements
- Crash-safe slave
- GTID-incompatible statement counters
- Source code available on GitHub: https://github.com/facebook/mysql-5.6/commits/webscalesql-5.6.16-47

### Why?

- No reasonable deployment plan exists in MySQL 5.6
- In MySQL documentation: Synchronize all the servers after stopping all the write traffic on master, restart all the servers in a replica cluster with gtid\_mode=ON. This is not a viable option at all
- Need a smooth rollout plan to reduce downtime of services

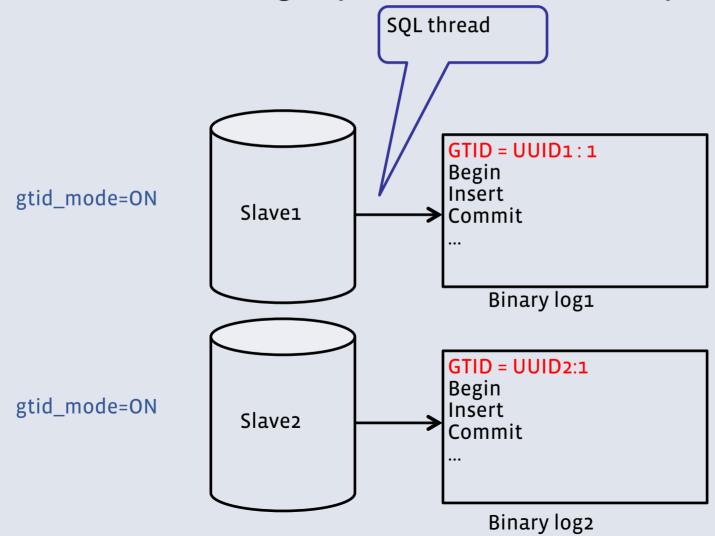


Try?



### Try?

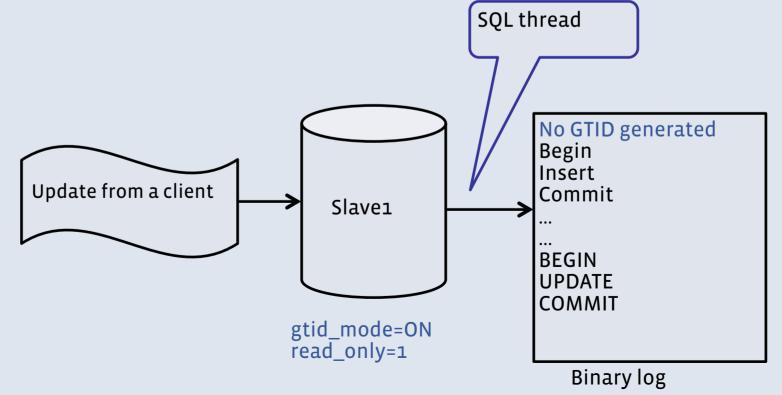
Restart slaves with --gtid\_mode=ON? Gtid\_executed will be different on each slave breaking replication with GTID protocol



#### How?

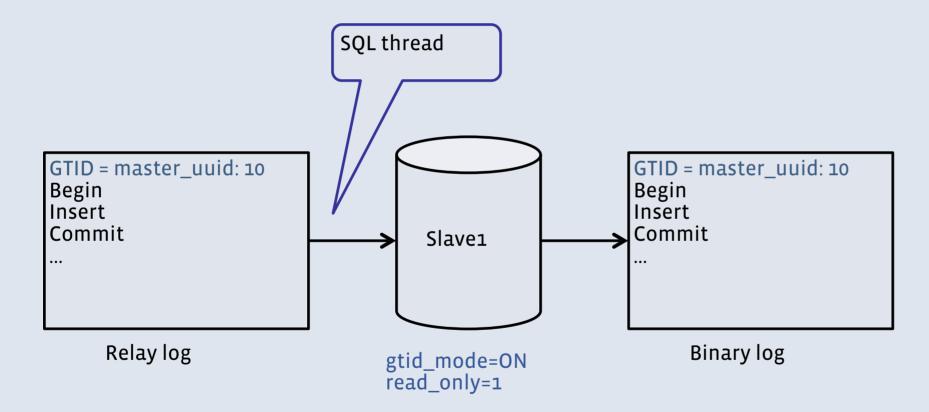
- Functionality to block generation of new GTIDs on slaves!
- Make sql\_thread never generate GTIDs, but log GTIDs from master
- Use read\_only=1 setting on slaves to block GTID generation

 Allow slave with -read\_only=1 to do replication from a master with gtid\_mode=OFF



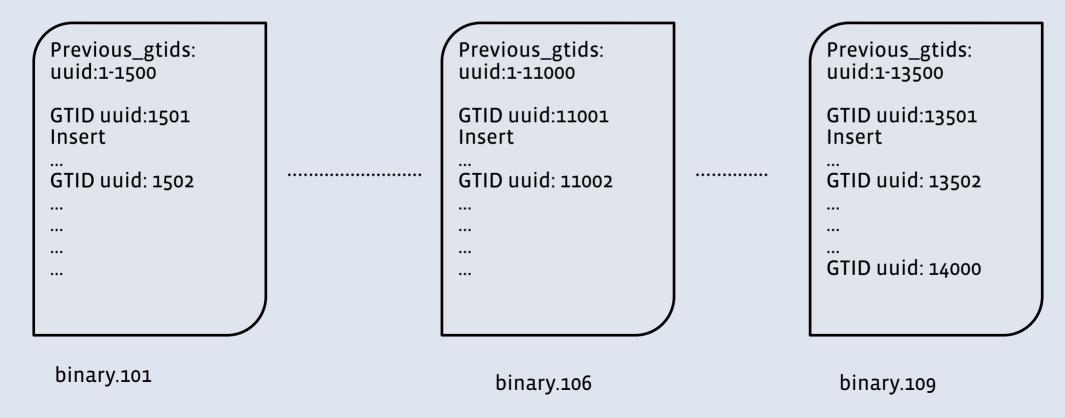
#### How?

SQL\_thread logs GTIDs received from master



# Slave connect performance

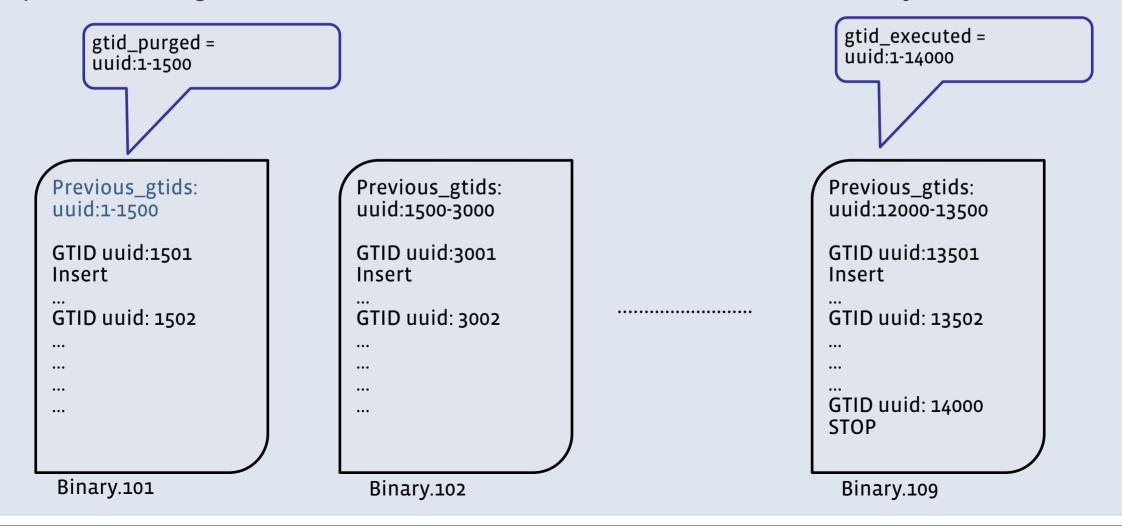
- In MySQL 5.6.10, dump thread scans all binary logs when slave connects with auto positioning
- After MySQL 5.6.10, dump thread opens binary logs in reverse order when slave connects with auto positioning. It opens logs from binary. 109 to binary.106 if slave's gtid\_executed='uuid:1-11000'



Binary logs on a master server

### Restart performance

- Mysqld opens binary logs in forward order on restart or crash recovery to initiate gtid\_set gtid\_purged and in reverse order to initiate gtid\_executed by using Previous\_gtid\_log\_events which are logged in the beginning of every binary log
- Opens all binlogs if GTIDs are not enabled, or GTIDs are enabled recently



# Change binary log index file format

- Keep the Previous\_gtid\_log\_events in memory and log them in the binary log index file along with the file name
- Use the previous\_gtid\_log\_events in memory to figure out the physical location when slave uses auto positioning
- On server restart, the Previous\_gtid\_log\_events in the binary log index file are used to initiate gtid\_sets

Slave uses BINLOG\_DUMP protocol instead of BINLOG\_DUMP\_GTID upon IO\_thread reconnect!

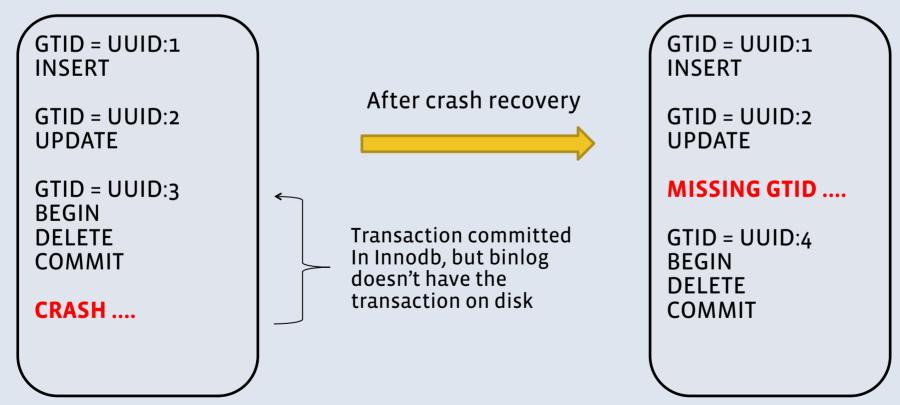
```
Binlog.100 (size of previous_gtid_log_event in binary format)
uuid:1-1500
binlog.00100
uuid:1500-3000
......
Binlog.106
uuid:1-11000
......
Binlog.109
uuid:12000-13500
```

### Crash-safe slave

- Running with full durability (sync\_binlog=1 and innodb\_flush\_log\_at\_trx\_commit =1) negatively affects performance
- We want slaves running with reduced durability (sync\_binlog != 1 and innodb\_flush\_log\_at\_trx\_commit != 1) to be consistent even after a soft crash or a hard crash
- MySQL 5.6 has crash safe functionality with the use of relay\_log\_info\_repository=TABLE. But it is not crash safe when GTIDs are enabled. Slave\_relay\_log\_info table only stores master binary log file name and position but doesn't have any GTID information

### Crash-safe slave: scenario 1

- If slave's binlog is behind Innodb transaction log during the crash, slave misses some GTIDs in binary logs that are committed in Innodb after crash recovery
- Slaves recovers if auto-positioning is OFF during the crash but will have missing GTIDs which will cause problems when auto-positioning is enabled later. But slave hits duplicate key error when auto-positioning is ON during the crash



### Crash-safe slave: scenario 1 fix

- Store the last committed GTID in a table slave\_gtid\_info. The transactional table is updated inside the SQL\_thread transaction or a slave worker in case of multi threaded slave
- Missing GTIDs in the binlog are identified using the slave\_gtid\_info table by checking the condition if current GTID of the SQL\_thread transaction is less than last committed GTID in the table. SQL\_thread logs those transactions directly into the binary log instead of executing them

```
SET @@GTID_NEXT = 'UUID:1332492';

START TRANSACTION;

INSERT INTO db1.tbl(col1, col2, blob)

UPDATE mysql.slave_gtid_info

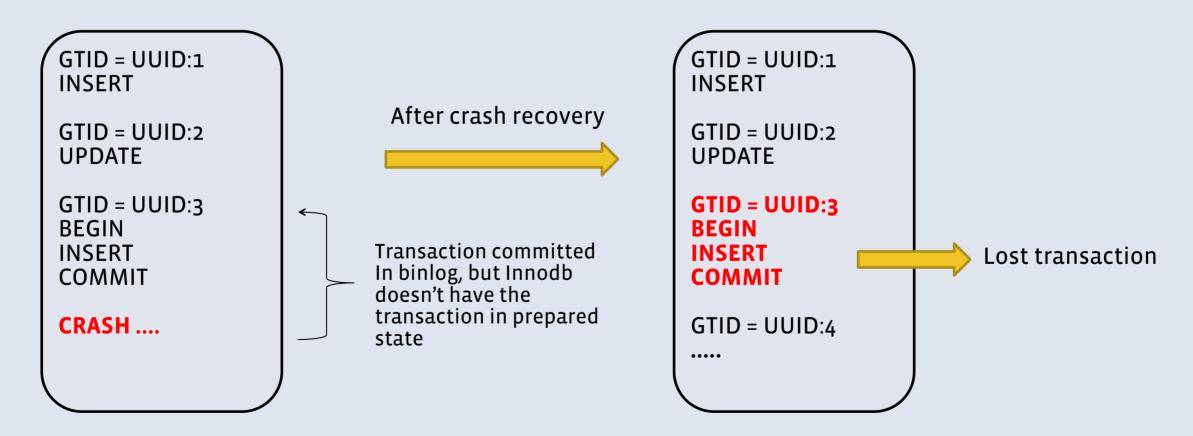
SET Last_gtid = 'UUID:1332492' where Database_name = 'db1'; // Added by SQL thread

COMMIT;
```

SQL\_thread executing a transaction

### Crash-safe slave: scenario 2

- If slaves binary log is ahead of Innodb transaction log during the crash, slave loses transactions
- Since slaves gtid\_executed contains UUID:3 after crash recovery, master skips sending the transaction corresponding to UUID:3



slaves binary log is ahead of Innodb transaction log

### Crash-safe slave: scenario 2 fix

 After crash recovery slave uses these physical binary log coordinates from Innodb log header to trim gtid\_executed before starting replication from master

### GTID bug reports

- http://bugs.mysql.com/bug.php?id=69059 GTIDs lack a reasonable deployment strategy
- http://bugs.mysql.com/bug.php?id=69097 MySQL scans all binary logs on crash recovery
- <a href="http://bugs.mysql.com/bug.php?id=68386">http://bugs.mysql.com/bug.php?id=68386</a> Master scans all binary logs when slave use auto-position
- http://bugs.mysql.com/bug.php?id=70659 Make GTID crash safe
- http://bugs.mysql.com/bug.php?id=71575 Master logs consecutive Gtid events
- http://bugs.mysql.com/bug.php?id=71695
   Concatenation of mysqlbinlog output does not work with GTID

## **Automation and Monitoring**

## Potential integration points

- Promotion logic
- Monitoring
- Config file generation
- Instance copying
- Backups (and restores!)
- Helper scripts for DBAs

#### **Promotion logic**

- Handle 3 cases of gtid\_mode:
  - Traditional: gtid\_mode=OFF on both old and new masters
  - GTID failover: gtid\_mode=ON on both old and new masters
  - Rollout: gtid\_mode=OFF on old master but ON for new master
- Many additional permutations to test
  - Which replica / binlog tailer is furthest ahead
  - Planned promotion (old master alive) vs failover (old master dead)

## Monitoring

- Broken replication from mismatched gtid\_mode (slave OFF, master ON)
- Replica sets using GTID but without auto-positioning not crash-safe
- Gaps in gtid\_executed
  - Should only occur if mysql server bug
  - Breaks auto-positioning
- Use of GTID-incompatible statements
  - Anything mixing DDL/nontransactional and transactional updates
  - FB-mysql adds user stat counters for these
- At Facebook, we use on-host agent plus external monitoring, and we also have external auto-remediation for solving common problems

## Config file generation

- gtid\_mode not dynamic in 5.6, so my.cnf is responsible for enabling
- Must support replica-set-at-a-time gtid\_mode rollout
- Programmatically set gtid\_mode and enforce\_gtid\_consistency
  - Enable if a flag has been set in asset tracker
  - Enable if gtid rollout complete for this entire logical data set
  - Enable if master already has gtid\_mode=ON
  - Disable if version < 5.6</li>

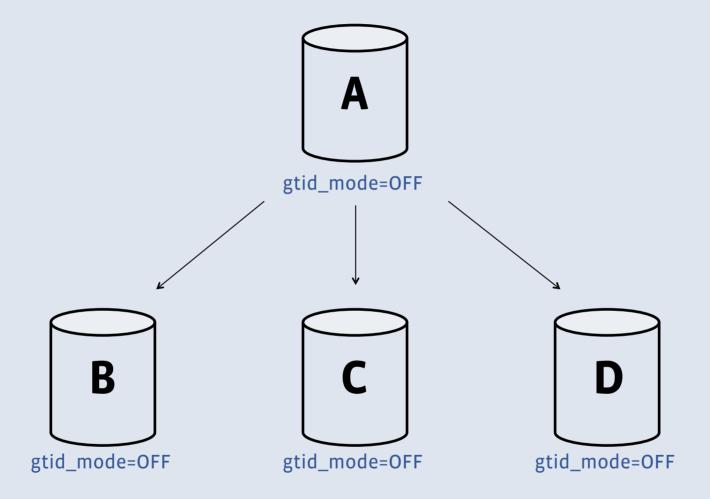
### Helper scripts for DBAs

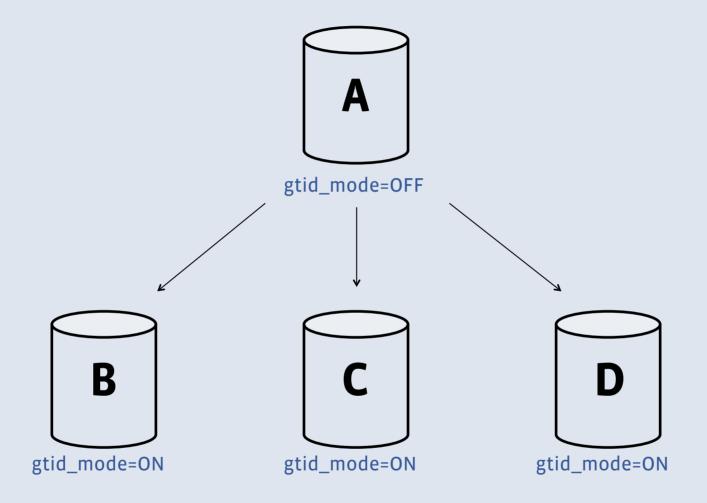
- Rollout script
- Rollback script? Rather painful
- Statement skipper
  - If gtid\_mode=OFF, use sql\_slave\_skip\_counter
  - If gtid\_mode=ON, need to insert empty transaction instead

## Deployment Strategy

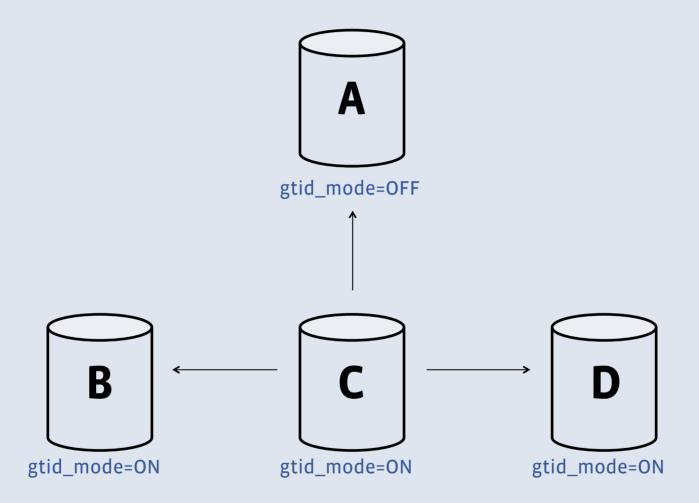
### Our rollout process

- 1. Set flag on replica set to enable gtid\_mode in my.cnf generator
- 2. Restart replicas one-at-a-time to enable gtid\_mode
  - Stop client conns and suppress monitoring/automation first
  - Prep fast shutdown by setting innodb\_max\_dirty\_pages\_pct=o
- 3. Perform a promotion
  - Brief splay time, to avoid too many concurrent promotions
  - New master has gtid\_mode=ON, so afterwards GTID fully in use
  - Still repoint old master, even though i/o thread will error upon start
- 4. Restart old master to enable gtid\_mode, replication will work now

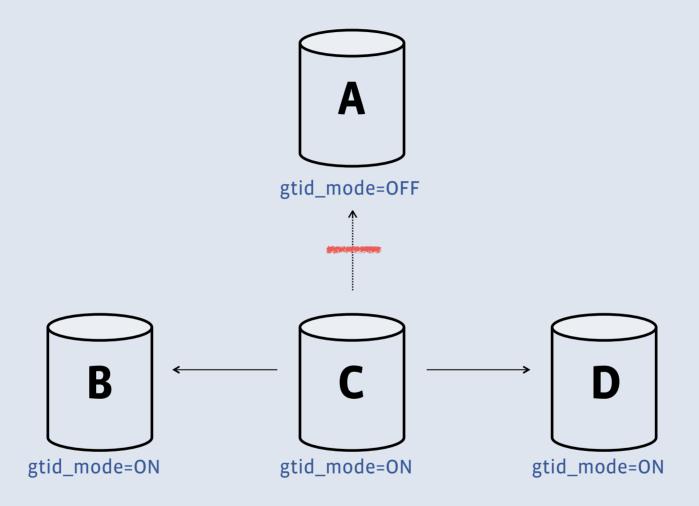




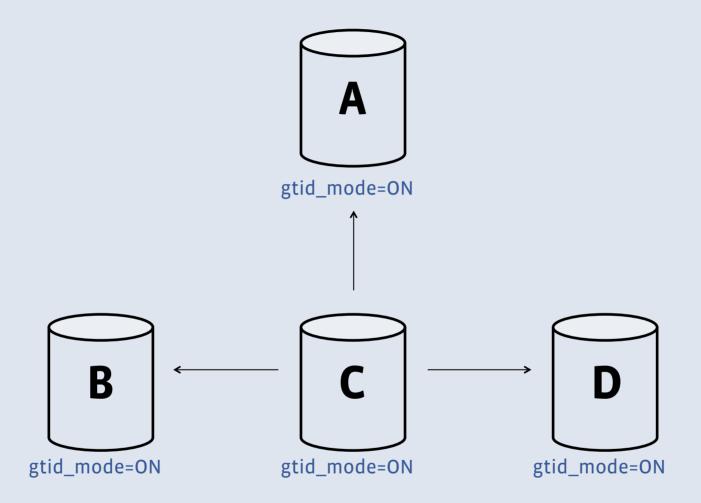
Enable gtid\_mode on replicas, one at a time



Promote C to be the new master



I/O thread on A breaks immediately due to gtid\_mode mismatch



Restart A to enable gtid\_mode and fix replication

#### Rollout validations

- Locking to prevent other automation from touching replica set. This includes copies, promotions, shard migrations, upgrades
- Confirm valid failover replica that isn't lagging
- User stat counters show no use of GTID-incompatible statements
- Replica set is reasonably healthy, no instances offline or broken
- All steps must succeed confirm clean shutdown and startup without hitting timeouts
- Validate gtid\_mode actually enabled after restart
- Confirm 5.6.x; if older 5.6.x, can upgrade instead of simple restart

## Questions?

# facebook

(c) 2007 Facebook, Inc. or its licensors. "Facebook" is a registered trademark of Facebook, Inc.. All rights reserved. 1.0